



UNIVERSITY OF PISA

Advanced Network Architectures and Wireless Systems

Year 2019/20

Centralized Virtual Redundancy Protocol

using a Floodlight OpenFlow controller

GERARDO ALVARO

LEONARDO FONTANELLI

RICCARDO POLINI

Introduction

The goal of this project is to define and implement a Floodlight Controller for SDN able to manage packets on a SDN switch (SS) and to route them from clients in Network A to servers in Network B through one gateway (either R1 or R2), as shown in Figure 1. Switch LS is considered as a legacy switch.

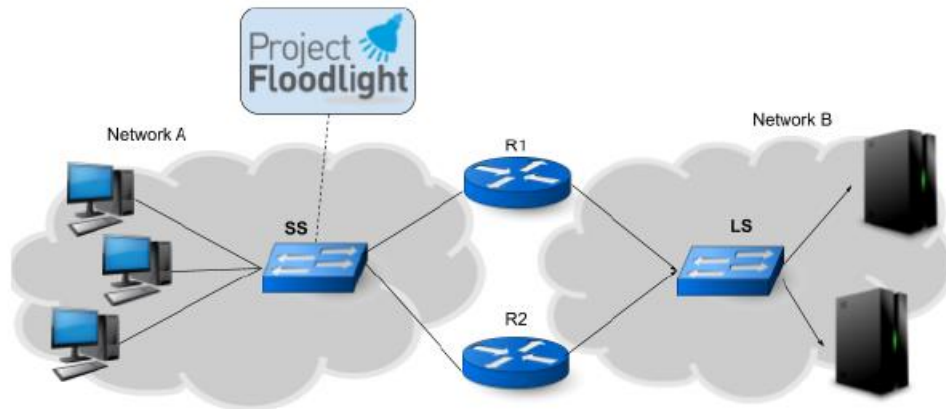


Figure 1: View of the global network

More specifically, at any time only one router operates as the gateway (Master), while the other one is available as a backup and will operate just in case of Master failure. Clients will be unaware of which router is acting as gateway. This will be achieved by masking the router's addresses with a virtual one.

Repository Structure

The project repository is organized as follows:

- **java:** contains 3 Java files (both controllers will be discussed later in detail):
 - **ARPController:** main file that handles ARP and ICMP packets.
 - **TController:** handles the router Redundancy protocol. Checks whether the master router is offline and, in case, elects a new master router.
 - **Parameters:** contains configurable parameters.
- **python:** contains the following Python files:
 - **net:** mininet script that builds the network configuration as shown in Figure 1.
 - **router:** contains the router operational code.

Redundancy Protocol

Guarantees the functionality of the network if the master router goes down (and a backup is available). Implemented in the **TController.java** file.

By default, and for a matter of simplicity, router R1 is the initial master and R2 is the backup. Each router sends a hello message consisting in the **ROUTER_ID**, every second on the UDP port 8787.

The **TController** module listens on port 8787 for router advertisements. When it receives one, it saves in a Hash Map the timestamp of arrival of the message, for that given router id.

The module also runs a timer which expires every 3 seconds. Every time an advertisement is received from the master router, the timer is reset. If the timer expires without having received an advertisement from the master router in the last 3 seconds, the router is considered offline and the backup router is elected as the new master (if available).

If no backup is available, the program is closed.

```
1586602949,63) R1 HELLO
1586602950,63) R1 HELLO
1586602951,63) R1 HELLO
1586602952,63) R1 HELLO
1586602953,63) R1 HELLO
1586602954,63) R1 HELLO
1586602955,63) R1 HELLO
1586602956,63) R1 HELLO
1586602957,64) R1 HELLO
1586602958,64) R1 HELLO
1586602959,64) R1 HELLO
1586602960,64) R1 HELLO
1586602961,64) R1 HELLO
1586602962,64) R1 HELLO
1586602963,64) R1 HELLO
1586602964) R1 HELLO
1586602965,65) R1 HELLO

Miracallback (most recent call last):
  File "router.py", line 44, in <module>
    advertise()
  File "router.py", line 40, in advertise
    time.sleep(T_ADV)
KeyboardInterrupt
root@riccardo-HP-Notebook:~/Scrivania/ANWS-Task2/python#

1586602955,21) R2 HELLO
1586602956,21) R2 HELLO
1586602957,21) R2 HELLO
1586602958,22) R2 HELLO
1586602959,22) R2 HELLO
1586602960,22) R2 HELLO
1586602961,22) R2 HELLO
1586602962,22) R2 HELLO
1586602963,22) R2 HELLO
1586602964,22) R2 HELLO
1586602965,22) R2 HELLO
1586602966,22) R2 HELLO
1586602967,23) R2 HELLO
1586602968,23) R2 HELLO
1586602969,23) R2 HELLO
1586602970,23) R2 HELLO
1586602971,23) R2 HELLO
1586602972,23) R2 HELLO
1586602973,23) R2 HELLO
1586602974,23) R2 HELLO
1586602975,24) R2 HELLO
1586602976,24) R2 HELLO
1586602977,24) R2 HELLO

[ava] Adding rule for send ICMP packets...
[ava] Adding rule for recv ICMP packets...
[ava] Processing IPv4 packet...
[ava] Adding rule for send ICMP packets...
[ava] Adding rule for recv ICMP packets...
[ava] Processing ARP request...
[ava] Managing incoming ARP Request from net B...
[ava] Processing ARP request...
[ava] Managing Virtual ARP Request...Sending out ARP reply
[ava] Processing ARP request...
[ava] Managing Virtual ARP Request...Sending out ARP reply
[ava] 2020-04-11 13:01:18.104 INFO [n.f.l.i.LinkDiscoveryManager] Sending LLDP packets out of all the enabled ports
[ava] 2020-04-11 13:01:33.111 INFO [n.f.l.i.LinkDiscoveryManager] Sending LLDP packets out of all the enabled ports
[ava] 2020-04-11 13:01:48.115 INFO [n.f.l.i.LinkDiscoveryManager] Sending LLDP packets out of all the enabled ports
[ava] 2020-04-11 13:02:03.120 INFO [n.f.l.i.LinkDiscoveryManager] Sending LLDP packets out of all the enabled ports
[ava] 2020-04-11 13:02:18.124 INFO [n.f.l.i.LinkDiscoveryManager] Sending LLDP packets out of all the enabled ports
[ava] 2020-04-11 13:02:33.130 INFO [n.f.l.i.LinkDiscoveryManager] Sending LLDP packets out of all the enabled ports
[ava] 2020-04-11 13:02:48.178 INFO [n.f.l.i.LinkDiscoveryManager] Sending LLDP packets out of all the enabled ports
[ava] Master Timer expired!
[ava] MASTER (1) is down. Attempting to connect to backup...
[ava] Router 2 is now MASTER.
[ava] Resetting flow rules...
```

In the above figure we can see that when R1 (master) is stopped, after the timer (3 seconds) expires, the backup router R2 is elected. If we restart R1, it will be set as backup router and will become master again only if R2 is stopped.

Address Plan

The following table shows the addresses we chose for the hosts of the network.

Host	IP
R1	10.0.2.1
R2	10.0.2.2
a-h1	10.0.2.3
a-h2	10.0.2.4
a-h3	10.0.2.5
b-h4	10.0.3.3
b-h5	10.0.3.4

As for the virtual addresses, we chose VIRTUAL_IP as **10.0.2.254** (last address of network A) and VIRTUAL_MAC as **00:00:5E:00:01:01** as suggested by RFC5798: *Virtual Router Redundancy Protocol (VRRP) Version 3 for IPv4 and IPv6*.

Virtual Address Masking

The n routers that lay in between network A and B are seen as one entity having a VIRTUAL_IP and a VIRTUAL_MAC. To implement this, the controller has to modify all packets involved with the routers. This means to replace the router's real addresses with the virtual ones.

When a packet arrives at the Floodlight controller, there can be two cases:

1. ARP packet

- a. **eth_source = REAL:** a host from network B is trying to communicate with network A. The gateway router sends an ARP Request to discover the address of the corresponding host in network A. The controller intercepts an ARP Request with the router's real MAC as *ethernet source* and sets the virtual addresses as sender addresses. A rule is saved in order to perform the same operation when the same conditions are intercepted again.
The controller also produces the inverse rule for the case when a host replies to the ARP Request. This rule consists in intercepting the host's ARP Reply and inserting the real router's addresses (host will reply to the virtual ones).
- b. **IP_destination = VIRTUAL:** the controller intercepts a broadcast ARP Request from network A which is trying to discover the VIRTUAL_MAC associated to the VIRTUAL_IP of the router. The controller builds an ARP Reply packet specifying the VIRTUAL_MAC and sends it back to the host.

2. ICMP packet

- a. **eth_destination = VIRTUAL:** the controller intercepts an ICMP packet from network A destined to the VIRTUAL_MAC and replaces it with the real router's MAC address. ICMP *source* and *destination* will still be host_A and host_B. A rule is saved in order to perform the same operation when the same conditions are intercepted again.
- b. **eth_source = REAL:** the controller intercepts an ICMP packet coming from router R2 (default gateway for network B). It sets VIRTUAL_MAC as the *ethernet source*.

Execution example

To test the software, the first thing to do is to launch the python script that generates the network by means of *mininet*.

```
sudo python net.py
```

Once inside the *mininet* console, we need to launch the two routers:

```
mininet> xterm r1 && xterm r2
```

Then we launch each router:

```
r1> python r1.py 1
```

```
r2> python r2.py 2
```

At this point we can launch floodlight and perform some tests.

ARP/ping example

We issued a ping request from h1 to h4 and tracked it with Wireshark on the interface between h1 and the switch.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	00:00:00_00:00:a1	Broadcast	ARP	42	Who has 10.0.2.254? Tell 10.0.2.3
2	0.006530052	IETF-VRRP-VRID_01	00:00:00_00:00:a1	ARP	42	10.0.2.254 is at 00:00:5e:00:01:01
3	0.006553922	10.0.2.3	10.0.3.3	ICMP	98	Echo (ping) request id=0x1aa4, seq=1/256, ttl=64 (reply in 6)
4	0.100785781	IETF-VRRP-VRID_01	Broadcast	ARP	42	Who has 10.0.2.3? Tell 10.0.2.254
5	0.100798148	00:00:00_00:00:a1	IETF-VRRP-VRID_01	ARP	42	10.0.2.3 is at 00:00:00:00:00:a1
6	0.101094205	10.0.3.3	10.0.2.3	ICMP	98	Echo (ping) reply id=0x1aa4, seq=1/256, ttl=63 (request in 3)
7	0.997243887	10.0.2.2	10.0.2.255	UDP	43	8787 → 8787 Len=1
8	1.001389347	10.0.2.3	10.0.3.3	ICMP	98	Echo (ping) request id=0x1aa4, seq=2/512, ttl=64 (reply in 9)
9	1.001900751	10.0.3.3	10.0.2.3	ICMP	98	Echo (ping) reply id=0x1aa4, seq=2/512, ttl=63 (request in 8)
10	1.998441835	10.0.2.2	10.0.2.255	UDP	43	8787 → 8787 Len=1
11	2.002468193	10.0.2.3	10.0.3.3	ICMP	98	Echo (ping) request id=0x1aa4, seq=3/768, ttl=64 (reply in 12)
12	2.002617014	10.0.3.3	10.0.2.3	ICMP	98	Echo (ping) reply id=0x1aa4, seq=3/768, ttl=63 (request in 11)
13	3.030401573	10.0.2.3	10.0.3.3	ICMP	98	Echo (ping) request id=0x1aa4, seq=4/1024, ttl=64 (reply in 14)
14	3.041746812	10.0.3.3	10.0.2.3	ICMP	98	Echo (ping) reply id=0x1aa4, seq=4/1024, ttl=63 (request in 13)

In message 1, h1 broadcasts an ARP Request to figure out the MAC corresponding to the **10.0.2.254** address (VIRTUAL IP). h1 receives an ARP Reply (produced by the controller) in which there is the VIRTUAL MAC (**00:00:5e:00:01:01**). At this point the ICMP requests/reply start working flawlessly.

Controller Rules example

What happens now on the controller side? We can open Wireshark on the Loopback interface and apply **tcp.port==6653** as a filter. if we run the previous example, we will see the FLOW_MOD of the controller.

tcp.port==6653					
No.	Time	Source	Destination	Protocol	Length Info
54911	13.986465229	127.0.0.1	127.0.0.1	OpenFlow	151 Type: OFPT_PACKET_IN
54912	13.986503374	127.0.0.1	127.0.0.1	TCP	66 6653 → 47050 [ACK] Seq=1829 Ack=2615
55615	14.183076243	127.0.0.1	127.0.0.1	OpenFlow	206 Type: OFPT_PACKET_IN
55616	14.183082594	127.0.0.1	127.0.0.1	TCP	66 6653 → 47050 [ACK] Seq=1829 Ack=2755
55617	14.183558631	127.0.0.1	127.0.0.1	OpenFlow	178 Type: OFPT_FLOW_MOD
55618	14.183565128	127.0.0.1	127.0.0.1	TCP	66 47050 → 6653 [ACK] Seq=2755 Ack=1941
55619	14.183787850	127.0.0.1	127.0.0.1	OpenFlow	178 Type: OFPT_FLOW_MOD
55620	14.183795036	127.0.0.1	127.0.0.1	TCP	66 47050 → 6653 [ACK] Seq=2755 Ack=2053
55621	14.183894019	127.0.0.1	127.0.0.1	OpenFlow	220 Type: OFPT_PACKET_OUT
55622	14.183899170	127.0.0.1	127.0.0.1	TCP	66 47050 → 6653 [ACK] Seq=2755 Ack=2207
56937	14.567179120	127.0.0.1	127.0.0.1	OpenFlow	178 Type: OFPT_PACKET_IN
56938	14.567877002	127.0.0.1	127.0.0.1	OpenFlow	176 Type: OFPT_PACKET_OUT
56939	14.567898954	127.0.0.1	127.0.0.1	TCP	66 47050 → 6653 [ACK] Seq=2867 Ack=2317
58494	14.987152043	127.0.0.1	127.0.0.1	OpenFlow	151 Type: OFPT_PACKET_IN
58669	15.030912401	127.0.0.1	127.0.0.1	TCP	66 6653 → 47050 [ACK] Seq=2317 Ack=2952
62210	15.988228317	127.0.0.1	127.0.0.1	OpenFlow	151 Type: OFPT_PACKET_IN
62211	15.988236429	127.0.0.1	127.0.0.1	TCP	66 6653 → 47050 [ACK] Seq=2317 Ack=3037
63054	16.199192414	127.0.0.1	127.0.0.1	OpenFlow	206 Type: OFPT_PACKET_IN
63055	16.199220045	127.0.0.1	127.0.0.1	TCP	66 6653 → 47050 [ACK] Seq=2317 Ack=3177
63060	16.200389002	127.0.0.1	127.0.0.1	OpenFlow	178 Type: OFPT_FLOW_MOD
63061	16.200411149	127.0.0.1	127.0.0.1	TCP	66 47050 → 6653 [ACK] Seq=3177 Ack=2429
63062	16.200750273	127.0.0.1	127.0.0.1	OpenFlow	178 Type: OFPT_FLOW_MOD
63063	16.200764170	127.0.0.1	127.0.0.1	TCP	66 47050 → 6653 [ACK] Seq=3177 Ack=2541
63064	16.200959856	127.0.0.1	127.0.0.1	OpenFlow	220 Type: OFPT_PACKET_OUT
63065	16.200971151	127.0.0.1	127.0.0.1	TCP	66 47050 → 6653 [ACK] Seq=3177 Ack=2695
63066	16.201490802	127.0.0.1	127.0.0.1	OpenFlow	206 Type: OFPT_PACKET_IN
63073	16.202518221	127.0.0.1	127.0.0.1	OpenFlow	178 Type: OFPT_FLOW_MOD
63074	16.202715110	127.0.0.1	127.0.0.1	OpenFlow	178 Type: OFPT_FLOW_MOD
63075	16.202860372	127.0.0.1	127.0.0.1	OpenFlow	220 Type: OFPT_PACKET_OUT
63076	16.202992189	127.0.0.1	127.0.0.1	TCP	66 47050 → 6653 [ACK] Seq=3317 Ack=3073
66041	16.909502061	127.0.0.1	127.0.0.1	OpenFlow	151 Type: OFPT_PACKET_IN
66198	17.030909313	127.0.0.1	127.0.0.1	TCP	66 6653 → 47050 [ACK] Seq=3073 Ack=3402
67339	17.383411044	127.0.0.1	127.0.0.1	OpenFlow	150 Type: OFPT_PACKET_IN
67340	17.383431597	127.0.0.1	127.0.0.1	TCP	66 6653 → 47050 [ACK] Seq=3073 Ack=3486
67341	17.383507585	127.0.0.1	127.0.0.1	OpenFlow	150 Type: OFPT_PACKET_IN
67342	17.383518896	127.0.0.1	127.0.0.1	TCP	66 6653 → 47050 [ACK] Seq=3073 Ack=3570

After each PACKET_IN there are 2 FLOW_MODs because we install a rule plus the reverse of that rule to allow for the response to flow correctly through the network.

In the image below we can see the contents of a FLOW_MOD packet. in this specific case the Match condition recognizes the VIRTUAL_MAC as the Ethernet Destination of a packet and installs the Action which consists in replacing the VIRTUAL_MAC with the current master router's MAC (in this case R1).

503374	127.0.0.1	127.0.0.1	TCP	66 6653
076243	127.0.0.1	127.0.0.1	OpenFlow	206 Type
082594				
558631				
565128				
787850				
795036				
894019				
899170				
179120				
877002				
898954				
152043				
912401				
228317				
236429				
192414				
220045				
389002				
411149				
750273				
764170				
959856				
971151				
490802				
518221				
715110				

Match

OXM field
 Class: OFPXM OFPXM OFPXM BASIC (0x8000)
 0000 011. = Field: OFPXM OFB_ETH_DST (3)
 0 = Has mask: False
 Length: 6
 Value: IETF-VRRP-VRID_01 (00:00:5e:00:01:01)

OXM field
 Pad: 00000000

Instruction
 Type: OFPIT_APPLY_ACTIONS (4)
 Length: 40
 Pad: 00000000

Action
 Type: OFPAT_SET_FIELD (25)
 Length: 16
 OXM field
 Class: OFPXM OFPXM OFPXM BASIC (0x8000)
 0000 011. = Field: OFPXM OFB_ETH_DST (3)
 0 = Has mask: False
 Length: 6
 Value: 00:00:00_00:00:01 (00:00:00:00:00:01)
 Pad: 0000

Action